

PRML (Pattern Recognition And Machine Learning) 读书会

第三章 Linear Models for Regression

主讲人 planktonli

QQ 群 177217565

读书会微信公众平台请扫描下面的二维码



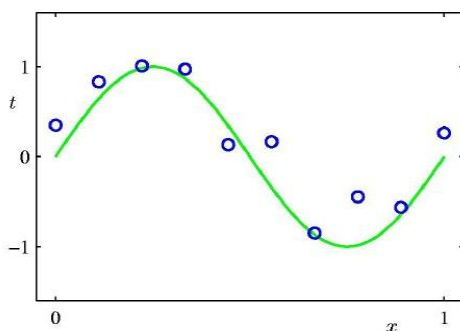
planktonli(1027753147) 18:58:12

大家好，我负责给大家讲讲 PRML 的第 3 讲 linear regression 的内容，请大家多多指教，群主让我们每个主讲人介绍下自己，赫赫，我也说两句，我是 applied mathematics + computer science 的，有问题大家可以直接指出，互相学习。大家有兴趣的话可以看看我的博客: <http://t.qq.com/keepuphero/mine>，当然我给大家推荐一个好朋友的，他对计算机发展还是很有心得的，他的网页 <http://www.zhizhihu.com/> 对 machine learning 的东西有深刻的了解。

好，下面言归正传，开讲第 3 章，第 3 章的名字是 linear regression，首先需要考虑的是：为什么在讲完 introduction、probability distributions 之后就直讲 linear regression? machine learning 的 essence 是什么？

机器学习的本质问题：我个人理解，就是通过数据集学习未知的最佳逼近函数，学习的收敛性\界 等等都是描述这个学习到的 function 到底它的性能如何。但是，从数学角度出发，函数是多样的，线性\非线性\跳跃\连续\非光滑，你可以组合出无数的函数，那么这些函数就组成了函数空间，在这些函数中寻找到一个满足你要求的最佳逼近函数，无疑大海捞针。我们再来回顾下第一章的 曲线拟和问题：

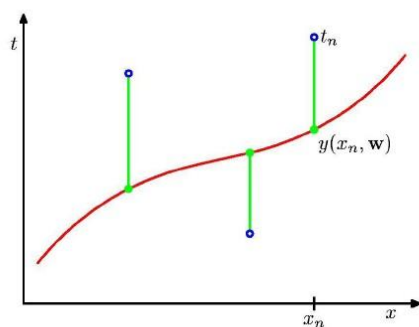
Polynomial Curve Fitting



$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

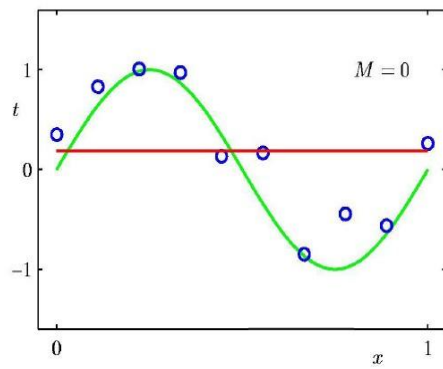
需要逼近的函数是: $\sin(2\pi x)$ ，M 阶的曲线函数可以逼近么？这是我们值得思考的问题。

Sum-of-Squares Error Function

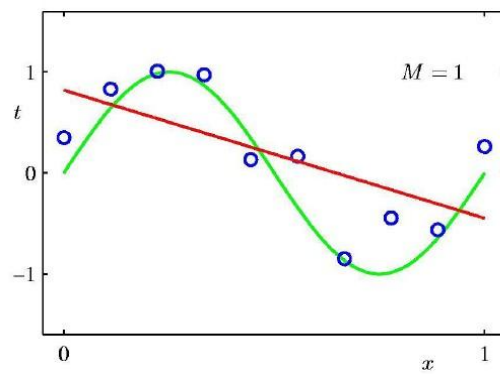


$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

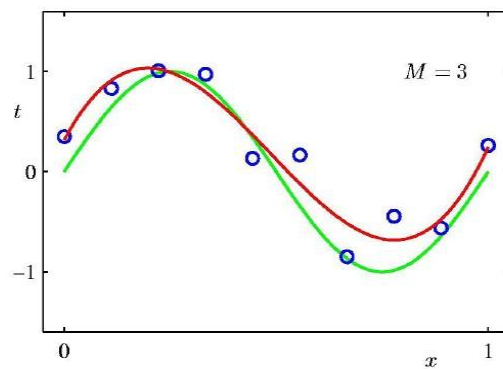
0th Order Polynomial



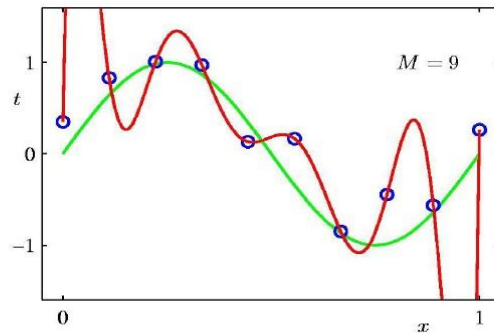
1st Order Polynomial



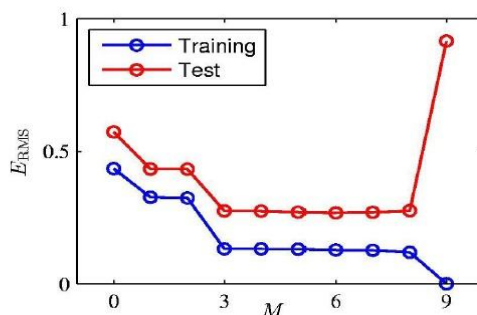
3rd Order Polynomial



9th Order Polynomial



Over-fitting



Root-Mean-Square (RMS) Error: $E_{\text{RMS}} = \sqrt{2E(\mathbf{w}^*)/N}$

要曲线拟和, 那么拟和的标准是什么?这里用了 2 范数定义,也就是误差的欧式距离, 当然,你可以用 L1,L 无穷, 等等了, 只是 objective 不同罢了。现在的疑问是: 为什么要用 Polynomial Fitting?有数学依据么, 这里牵扯到 范函的问题, 就是函数所张成的空间, 举一个简单的例子, 大家还都记得 taylor 展式吧:

$$f(x) = f(x_0) + \frac{f'(x)}{1!}(x - x_0) + \frac{f''(x)}{2!}(x - x_0)^2 + \dots$$

这表明 任意一个函数可以表示成 x 的次方之和, 也就是 任意一个函数 可以放到 $(1, x, x^2, x^3, \dots)$ 所张成的函数空间, 如果是有限个基的话就称为欧式空间, 无穷的话 就是 Hilbert 空间, 其实 傅里叶变换 也是这样的例子, 既然已经明白了 任意函数可以用 Polynomial Fitting, 那么下面就是什么样的 Polynomial 是最好的。

Wilbur_中博(1954123) 19:28:26

泰勒展开是局部的、 x_0 周围的, 而函数拟合是全局的, 似乎不太一样吧?

planktonli(1027753147) 19:29:21

恩,泰勒展开是局部的, 他是在 x_0 点周围的一个 表达, 函数拟合是全局的, 我这里只是用一个简单的例子说明 函数表达的问题。

Wilbur_中博(1954123) 19:30:41



planktonli(1027753147) 19:31:03

其实,要真正解释这个问题是需要范函的东西的。

Wilbur_中博(1954123) 19:31:45

抱歉，打断了一下，因为我觉得这个问题留到讨论就不太好了，呵呵。了解了，请继续吧。

planktonli(1027753147) 19:31:51

由于大多数群友未学过这个课程,我只是想说下这个思想,呵呵,没事,讨论才能深刻理解问题,其实,wavelet 这些,包括 kernel construcion 这些东西都牵扯到 范函。

Bishop 用上面这个例子说明：

1) 可以用 Polynomial Fitting 拟和 sin 类的函数 2) 存在过拟和问题

而且这里的 Polynomial Fitting 是一个线性 model，这里 Model 是 w 的函数, w 是线性的：

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

$\sin(2\pi x)$ 是线性的么，肯定不是，那么 让我们再分析下 研究的问题

$\sin(2\pi x)$ 中的 x 是 1 维的

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

上面的 x 变成了 $[1, x, x^2, x^3, \dots]^T$

$y(x, w) = W^T X$ ，非常有意思的是：维数升高了，同时这个 model 具有了表达非线性东西的能力。这

里的思想,可以说贯穿在 NN,SVM 这些东西里，也就是说,线性的 model 如果应用得当的话,可以表达非线性的东西。与其在所有函数空间盲目的寻找,还不如从一个可行的简单 model 开始，这就是为什么 Bishop 在讲完基础后直接切入 Linear regression 的原因,当然这个线性 model 怎么构造,是单层的 linear model,还是多层的 linear model 一直争论不休，BP 否定了 perceptron 的 model，SVM 否定了 BP model 现在 deep learning 又质疑 SVM 的 shallow model，或许这就是 machine learning 还能前进的动力。

让咱们再回来看看 linear regression 的模型，这里从标准形式到扩展形式，也就是引入基函数后,Linear regression 的模型可以表达非线性的东西了，因为基函数可能是非线性的：

Linear Regression

The basic form of linear regression

$$y(\mathbf{x}, \mathbf{w}) = w_0 + w_1x_1 + \dots + w_Dx_D$$

where $\mathbf{x} = (x_1, \dots, x_D)^T$ is the input variable, and $\mathbf{w} = (w_1, \dots, w_D)^T$ is the parameters.

More general form of linear regression

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j\phi_j(\mathbf{x})$$

where $\phi_j(\mathbf{x})$'s are known as *basis functions*.

Parameter w_0 is called a *bias* parameter. By adding $\phi_0(\mathbf{x}) = 1$, we have

$$y(\mathbf{x}, \mathbf{w}) = \sum_{j=0}^{M-1} w_j\phi_j(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x})$$

基函数的形式，这些基函数都是非线性的：

Basis functions

Polynomial functions

$$\phi_j(x) = x^j$$

Gaussian functions

$$\phi_j(\mathbf{x}) = \exp \left\{ -\frac{(\mathbf{x} - \boldsymbol{\mu}_j)^T (\mathbf{x} - \boldsymbol{\mu}_j)}{2s^2} \right\}$$

Sigmoid basis functions

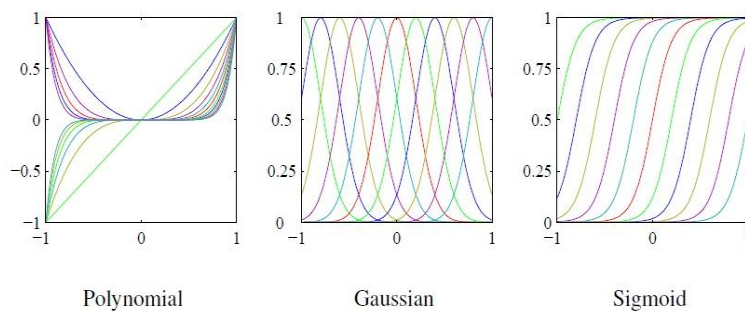
$$\phi_j(\mathbf{x}) = \sigma \left(\frac{\mathbf{b}^T \mathbf{x} - \mu_j}{s} \right)$$

where $\sigma(a)$ is the logistic sigmoid function

$$\sigma(a) = \frac{1}{1 + \exp \{-a\}}$$

Other basis functions: wavelets

Basis functions



Probabilistic Formulation

Assume the target variable t is given by a deterministic function $y(\mathbf{x}, \mathbf{w})$ with additive Gaussian noise

$$t = y(\mathbf{x}, \mathbf{w}) + \epsilon$$

where ϵ is a zero mean Gaussian random variable with precision (inverse variance) β .
Therefore

$$p(t|\mathbf{x}, \mathbf{w}, \beta) = \mathcal{N}(t|y(\mathbf{x}, \mathbf{w}), \beta^{-1})$$

Conditional mean

$$\mathbb{E}[t|\mathbf{x}] = \int t p(t|\mathbf{x}) dt = y(\mathbf{x}, \mathbf{w}) + \mathbb{E}[\epsilon] = y(\mathbf{x}, \mathbf{w})$$

在 Gaussian 零均值情况下, Linear model 从频率主义出发的 MLE 就是 Least square :

Maximum Likelihood Estimator

Consider a data set of inputs $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ with corresponding target values $\mathbf{t} = (t_1, \dots, t_N)^T$. The likelihood function (a function of \mathbf{w} and β) is

$$p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \prod \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1})$$

Log likelihood

$$\begin{aligned} \log p(\mathbf{t}|\mathbf{w}, \beta) &= \sum_{n=1}^N \log \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1}) \\ &= \frac{N}{2} \log \beta - \frac{N}{2} \log 2\pi - \beta E_D(\mathbf{w}) \end{aligned}$$

where the sum-of-square error function is

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (t_n - \mathbf{w}^T \phi(\mathbf{x}_n))^2$$

最小 2 乘的解就是广义逆矩阵乘输出值 :

Maximum Likelihood Estimator: \mathbf{w}

With β fixed, we obtain the gradient of the log likelihood

$$\nabla \log p(\mathbf{t}|\mathbf{w}, \beta) = \sum_{n=1}^N (t_n - \mathbf{w}^T \phi(\mathbf{x}_n)) \phi(\mathbf{x}_n)^T$$

Setting the gradient to be zero gives

$$\begin{aligned} \mathbf{w}^T \left(\sum_{n=1}^N \phi(\mathbf{x}_n) \phi(\mathbf{x}_n)^T \right) &= \sum_{n=1}^N t_n \phi(\mathbf{x}_n)^T \\ \mathbf{w}^T \Phi^T \Phi &= \mathbf{t}^T \Phi \\ \Phi^T \Phi \mathbf{w} &= \Phi^T \mathbf{t} \\ \mathbf{w}_{\text{ML}} &= (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t} \end{aligned}$$

where the $N \times M$ design matrix Φ is $\Phi = [\phi(\mathbf{x}_1) \dots \phi(\mathbf{x}_N)]^T$. The quantity $\Phi^\dagger \equiv (\Phi^T \Phi)^{-1} \Phi^T$ is called *pseudo-inverse* of the matrix Φ .

Gaussian 的 precision 也可以计算出来 :

Maximum Likelihood Estimator: β

Recall log likelihood

$$\log p(\mathbf{t}|\mathbf{w}, \beta) = \frac{N}{2} \log \beta - \frac{N}{2} \log 2\pi - \beta E_D(\mathbf{w})$$

Setting the partial derivative w.r.t. β to zero

$$\frac{\partial}{\partial \beta} \log p(\mathbf{t}|\mathbf{w}, \beta) = \frac{N}{2} \frac{1}{\beta} - E_D(\mathbf{w})$$

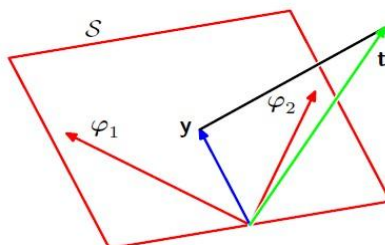
We obtain

$$\beta_{\text{ML}} = \frac{N}{(\mathbf{t} - \Phi \mathbf{w}_{\text{ML}})^T (\mathbf{t} - \Phi \mathbf{w}_{\text{ML}})}$$

最小 2 乘的解可以看成到基张成空间的投影：

Geometrical Interpretation of Least Squares

The least-squares regression function is obtained by finding the orthogonal projection of the data vector \mathbf{t} onto the subspace spanned by the basis functions $\phi_j(\mathbf{x})$ in which each basis function is viewed as a vector φ_j of length $|N|$ with elements $\phi_j(\mathbf{x}_n)$.



频率主义会导致 过拟和，加入正则,得到的最小 2 乘解：

Regularized Least Squares

The over-fitting issue can be addressed through adding a *regularization* term

$$E_D(\mathbf{w}) + \lambda E_W(\mathbf{w})$$

where λ is the regularization coefficient. For example

$$E_W(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

and the error function becomes

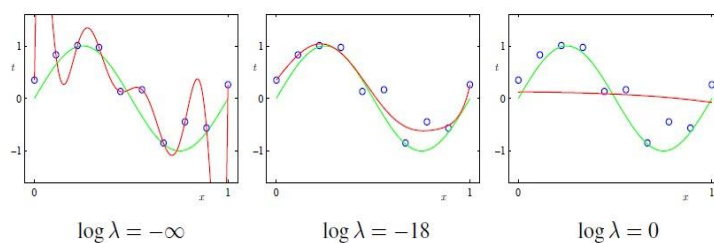
$$\frac{1}{2} (\mathbf{t} - \Phi \mathbf{w})^T (\mathbf{t} - \Phi \mathbf{w}) + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

which leads to solution

$$\mathbf{w} = (\lambda \mathbf{I} + \Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

正则参数对 model 结果的影响：

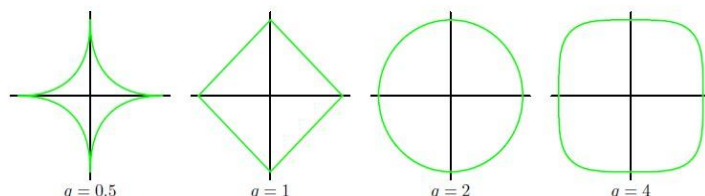
Effects of λ



消除过拟和，正则的几何解释：

Other Regularizers

$$E_W(\mathbf{w}) = \frac{1}{2} \sum_{j=1}^M |w_j|^q$$



When $q = 1$, the regularizer is called *lasso* in statistics as sparse models are preferred.

正则方法不同,就会出现很多 model,例如 lasso, ridge regression。LASSO 的解是稀疏的,例如:sparse coding,Compressed sensing 是从 L0--> L1sparse 的问题,现在也很热的。

Regularized Least Squares and MAP



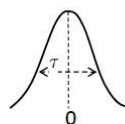
What if $(\mathbf{A}^T \mathbf{A})$ is not invertible ?

$$\hat{\beta}_{\text{MAP}} = \arg \max_{\beta} \underbrace{\log p(\{(X_i, Y_i)\}_{i=1}^n | \beta, \sigma^2)}_{\text{log likelihood}} + \underbrace{\log p(\beta)}_{\text{log prior}}$$

I) Gaussian Prior

$$\beta \sim \mathcal{N}(0, \tau^2 \mathbf{I})$$

$$p(\beta) \propto e^{-\beta^T \beta / 2\tau^2}$$



$$\hat{\beta}_{\text{MAP}} = \arg \min_{\beta} \sum_{i=1}^n (Y_i - X_i \beta)^2 + \lambda \|\beta\|_2^2 \quad \text{Ridge Regression}$$

Closed form: HW constant(σ^2, τ^2)

Regularized Least Squares and MAP



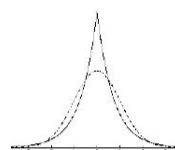
What if $(\mathbf{A}^T \mathbf{A})$ is not invertible ?

$$\hat{\beta}_{\text{MAP}} = \arg \max_{\beta} \underbrace{\log p(\{(X_i, Y_i)\}_{i=1}^n | \beta, \sigma^2)}_{\text{log likelihood}} + \underbrace{\log p(\beta)}_{\text{log prior}}$$

II) Laplace Prior

$$\beta_i \stackrel{iid}{\sim} \text{Laplace}(0, t)$$

$$p(\beta_i) \propto e^{-|\beta_i|/t}$$



$$\hat{\beta}_{\text{MAP}} = \arg \min_{\beta} \sum_{i=1}^n (Y_i - X_i \beta)^2 + \lambda \|\beta\|_1 \quad \text{Lasso}$$

Closed form: HW constant(σ^2, t)

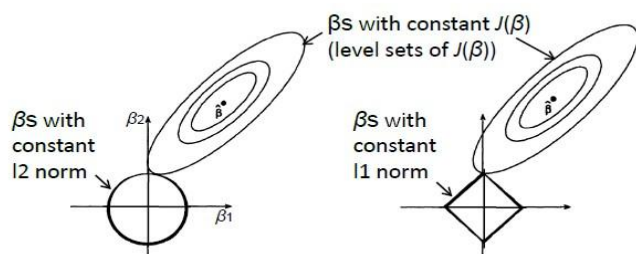
Ridge Regression vs Lasso

$$\min_{\beta} (\mathbf{A}\beta - \mathbf{Y})^T (\mathbf{A}\beta - \mathbf{Y}) + \lambda \text{pen}(\beta) = \min_{\beta} J(\beta) + \lambda \text{pen}(\beta)$$

Ridge Regression:
 $\text{pen}(\beta) = \|\beta\|_2^2$

Lasso:
 $\text{pen}(\beta) = \|\beta\|_1$

HOT!



Lasso (l1 penalty) results in sparse solutions – vector with more zero coordinates
Good for high-dimensional problems – don't have to store all coordinates!

下面看 Bias-Variance Decoposition , 正则就是在 训练数据的模型上加一个惩罚项, shrink 模型的参数,

让它不要学习的太过, 这里 $E_D(\mathbf{w})$ 是对训练数据学习到的模型, $E_W(\mathbf{w})$ 是学习到的参数的惩罚模型

Expected Squared Loss

Recall that t is the value for each input \mathbf{x} and $y(\mathbf{x})$ is our estimate. We incur a loss $L(t, y(\mathbf{x}))$, and the expected loss is

$$\mathbb{E}[L] = \int \int L(t, y(\mathbf{x})) p(\mathbf{x}, t) d\mathbf{x} dt.$$

A common choice of loss function is the squared loss $L(t, y(\mathbf{x})) = \{y(\mathbf{x}) - t\}^2$, and the expected squared loss is

$$\mathbb{E}[L] = \int \int \{y(\mathbf{x}) - t\}^2 p(\mathbf{x}, t) d\mathbf{x} dt$$

Our goal is to choose $y(\mathbf{x})$ to minimize $\mathbb{E}[L]$. From Euler-Lagrange

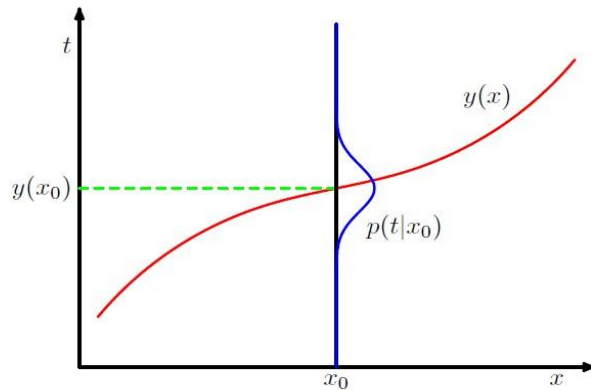
$$\frac{\delta \mathbb{E}[L]}{\delta y(\mathbf{x})} = 2 \int \{y(\mathbf{x}) - t\} p(\mathbf{x}, t) dt = 0$$

So

$$y(\mathbf{x}) = \frac{\int t p(\mathbf{x}, t) dt}{p(\mathbf{x})} = \int t p(t|\mathbf{x}) dt = \mathbb{E}_t[t|\mathbf{x}]$$

Regression Function

The function $y(\mathbf{x}) = \mathbb{E}[t|\mathbf{x}]$ is known as the *regression function*.



Decomposition of the Loss Function

Define

$$h(\mathbf{x}) = \mathbb{E}[t|\mathbf{x}] = \int t p(t|\mathbf{x}) dt$$

Since the optimal solution is the conditional expectation, we can expand the square term

$$\begin{aligned} \{y(\mathbf{x}) - t\}^2 &= \{y(\mathbf{x}) - \mathbb{E}[t|\mathbf{x}] + \mathbb{E}[t|\mathbf{x}] - t\}^2 \\ &= \{y(\mathbf{x}) - h(\mathbf{x})\}^2 + 2\{y(\mathbf{x}) - h(\mathbf{x})\}\{h(\mathbf{x}) - t\} + \{h(\mathbf{x}) - t\}^2 \end{aligned}$$

Therefore

$$\begin{aligned} \mathbb{E}[L] &= \int \int \{y(\mathbf{x}) - t\}^2 p(\mathbf{x}, t) d\mathbf{x} dt \\ &= \int \{y(\mathbf{x}) - h(\mathbf{x})\}^2 p(\mathbf{x}) d\mathbf{x} + \int \{h(\mathbf{x}) - t\}^2 p(\mathbf{x}, t) d\mathbf{x} dt \end{aligned}$$

When $y(\mathbf{x}) = h(\mathbf{x})$, the first term will vanish. The second term is the variance of the distribution of t averaged over \mathbf{x} , representing the intrinsic variability of the target data and can be regarded as noise and the irreducible minimum value of the loss function.

Decomposition of the Loss Function

In practice, $h(\mathbf{x})$ is unknown. We have a data set \mathcal{D} containing only a finite number N of data points. The loss function becomes

$$\{y(\mathbf{x}; \mathcal{D}) - h(\mathbf{x})\}^2$$

We can rewrite it as

$$\begin{aligned} \{y(\mathbf{x}; \mathcal{D}) - h(\mathbf{x})\}^2 &= \{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] + \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\}^2 \\ &= \{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})]\}^2 + \{\mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\}^2 \\ &\quad + 2\{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})]\}\{\mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\} \end{aligned}$$

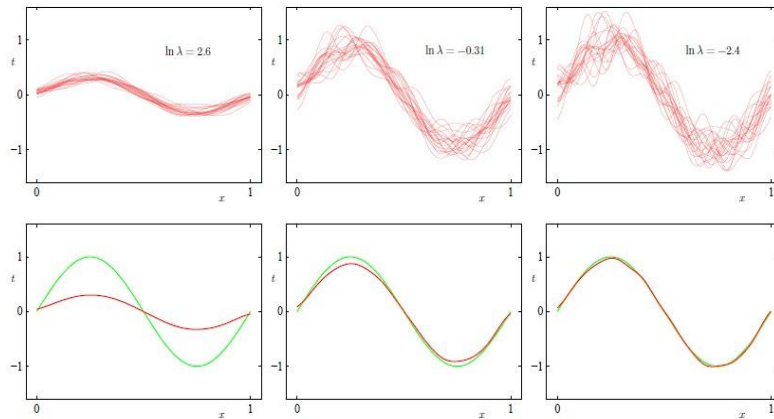
Take the expectation w.r.t. \mathcal{D}

$$\begin{aligned} \mathbb{E}_{\mathcal{D}}[\{y(\mathbf{x}; \mathcal{D}) - h(\mathbf{x})\}^2] &= \underbrace{\{\mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\}^2}_{\text{bias}^2} + \underbrace{\mathbb{E}_{\mathcal{D}}[\{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})]\}^2]}_{\text{variance}} \end{aligned}$$

上面这么多 PPT 无非就是说，学习到的模型和真实的模型的期望由 2 部分组成：

1--> Bias 2--> Variance。Bias 表示的是学习到的模型和真实模型的偏离程度,Variance 表示的是学习到的模型和它自己的期望的偏离程度。从这里可以看到正则项在控制 Bias 和 Variance :

Bias-Variance as a Function of Complexity



Wilbur_中博(1954123) 20:33:07

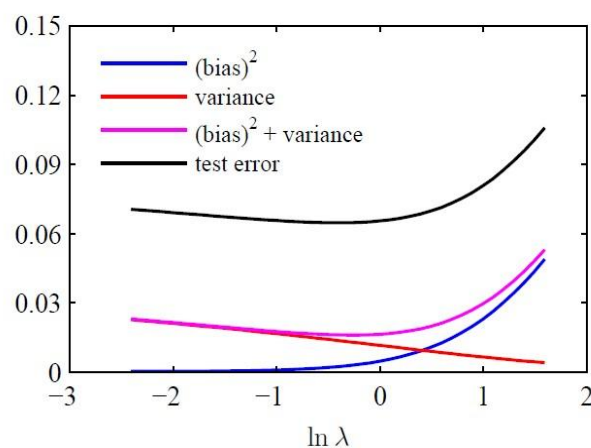
这个是关键，呵呵

planktonli(1027753147) 20:33:25

Variance 小的情况下,Bias 就大 , Variance 大的情况下,Bias 就小 , 我们就要 tradeoff 它们。

从这张图可以看到 Bias 和 Variance 的关系 :

Bias-Variance as a Function of Complexity



这个 Bias-Variance Decoposition 其实没有太大的实用价值，它只能起一个指导作用。

下面看看 Bayesian Linear Regression :

Conjugate Prior

Recall the likelihood function

$$\begin{aligned} p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) &= \prod_{n=1}^N \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1}) \\ &= \mathcal{N}(\mathbf{t} | \Phi \mathbf{w}, \beta^{-1} \mathbf{I}) \end{aligned}$$

is the exponential of a quadratic function of \mathbf{w} . So the conjugate prior is given by a Gaussian distribution of the form

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w} | \mathbf{m}_0, \mathbf{S}_0)$$

Bayesian Theorem for Gaussian

Theorem (Bayesian Theorem for Linear Gaussian)

Given prior and likelihood

$$\begin{aligned} p(\mathbf{x}) &= \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}, \boldsymbol{\Lambda}^{-1}) \\ p(\mathbf{y}|\mathbf{x}) &= \mathcal{N}(\mathbf{y} | \mathbf{A}\mathbf{x} + \mathbf{b}, \mathbf{L}^{-1}) \end{aligned}$$

the marginal $p(\mathbf{y})$ and posterior $p(\mathbf{x}|\mathbf{y})$ are

$$\begin{aligned} p(\mathbf{y}) &= \mathcal{N}(\mathbf{y} | \mathbf{A}\boldsymbol{\mu} + \mathbf{b}, \mathbf{A}\boldsymbol{\Lambda}^{-1}\mathbf{A}^T + \mathbf{L}^{-1}) \\ p(\mathbf{x}|\mathbf{y}) &= \mathcal{N}(\mathbf{x} | \boldsymbol{\Sigma}[\mathbf{A}^T\mathbf{L}(\mathbf{y} - \mathbf{b}) + \mathbf{A}\boldsymbol{\mu}], \boldsymbol{\Sigma}) \end{aligned}$$

where

$$\boldsymbol{\Sigma} = (\boldsymbol{\Lambda} + \mathbf{A}^T\mathbf{L}\mathbf{A})^{-1}.$$

Posterior

Apply the theorem and we obtain the posterior

$$p(\mathbf{w}|\mathbf{t}) = \mathcal{N}(\mathbf{w} | \mathbf{m}_N, \mathbf{S}_N)$$

where

$$\begin{aligned} \mathbf{m}_N &= \mathbf{S}_N(\mathbf{S}_0^{-1}\mathbf{m}_0 + \beta\Phi^T\mathbf{t}) \\ \mathbf{S}_N^{-1} &= \mathbf{S}_0^{-1} + \beta\Phi^T\Phi \end{aligned}$$

The prior can be a zero-mean isotropic Gaussian governed by a single precision parameter α

$$p(\mathbf{w}|\alpha) = \mathcal{N}(\mathbf{w} | \mathbf{0}, \alpha^{-1}\mathbf{I})$$

and the corresponding posterior distribution is

$$\begin{aligned} \mathbf{m}_N &= \beta\mathbf{S}_N\Phi^T\mathbf{t} \\ \mathbf{S}_N^{-1} &= \alpha\mathbf{I} + \beta\Phi^T\Phi \end{aligned}$$

Bayesian MAP vs. Regularization

The log posterior takes the following form

$$\log p(\mathbf{w}|\mathbf{t}) = -\frac{\beta}{2}(\mathbf{t} - \Phi\mathbf{w})^T(\mathbf{t} - \Phi\mathbf{w}) - \frac{\alpha}{2}\mathbf{w}^T\mathbf{w} + \text{const}$$

The regularized least square takes the form

$$\frac{1}{2}(\mathbf{t} - \Phi\mathbf{w})^T(\mathbf{t} - \Phi\mathbf{w}) + \frac{\lambda}{2}\mathbf{w}^T\mathbf{w}$$

Therefore, $\lambda = \frac{\alpha}{\beta}$.

从 Bayesian 出发,关注的不是参数的获取,而更多的是 新预测的值,通过后验均值可以得到 linear model 和核函数的联系,当然也可以建立 gaussian process 这些东西。

Wilbur_中博(1954123) 20:51:25

这里可以讲细一点么,如何建立联系?

planktonli(1027753147) 20:54:44

Bayesian Linear Regression (2)

A common choice for the prior is

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \alpha^{-1}\mathbf{I})$$

for which

$$\begin{aligned}\mathbf{m}_N &= \beta\mathbf{S}_N\Phi^T\mathbf{t} \\ \mathbf{S}_N^{-1} &= \alpha\mathbf{I} + \beta\Phi^T\Phi.\end{aligned}$$

Equivalent Kernel (1)

The predictive mean can be written

$$\begin{aligned}y(\mathbf{x}, \mathbf{m}_N) &= \mathbf{m}_N^T\phi(\mathbf{x}) = \beta\phi(\mathbf{x})^T\mathbf{S}_N\Phi^T\mathbf{t} \\ &= \sum_{n=1}^N \underbrace{\beta\phi(\mathbf{x})^T\mathbf{S}_N\phi(\mathbf{x}_n)}_{k(\mathbf{x}, \mathbf{x}_n)} t_n \\ &= \sum_{n=1}^N k(\mathbf{x}, \mathbf{x}_n) t_n.\end{aligned}$$

Equivalent kernel or smoother matrix.

This is a weighted sum of the training data target values, t_n .

这里就可以看到了啊,看到了么, Wilbur?

Wilbur_中博(1954123) 20:57:24

👉 在看

planktonli(1027753147) 20:58:08

如果共扼先验是 0 均值情况下,linear model 就可以变成 kernel 了:

Equivalent Kernel (4)

The kernel as a covariance function: consider

$$\begin{aligned}\text{cov}[y(\mathbf{x}), y(\mathbf{x}')] &= \text{cov}[\phi(\mathbf{x})^T \mathbf{w}, \mathbf{w}^T \phi(\mathbf{x}')] \\ &= \phi(\mathbf{x})^T \mathbf{S}_N \phi(\mathbf{x}') = \beta^{-1} k(\mathbf{x}, \mathbf{x}').\end{aligned}$$

We can avoid the use of basis functions and define the kernel function directly, leading to *Gaussian Processes* (Chapter 6).

Equivalent Kernel (5)

$$\sum_{n=1}^N k(\mathbf{x}, \mathbf{x}_n) = 1$$

for all values of \mathbf{x} ; however, the equivalent kernel may be negative for some values of \mathbf{x} .

Like all kernel functions, the equivalent kernel can be expressed as an inner product:

$$k(\mathbf{x}, \mathbf{z}) = \psi(\mathbf{x})^T \psi(\mathbf{z})$$

where $\psi(\mathbf{x}) = \beta^{1/2} \mathbf{S}_N^{1/2} \phi(\mathbf{x})$.

最后讲了 bayesian model 比较 :

Bayesian Model Comparison (1)

How do we choose the 'right' model?

Assume we want to compare models \mathcal{M}_i , $i=1, \dots, L$, using data \mathcal{D} ; this requires computing

$$p(\mathcal{M}_i | \mathcal{D}) \propto p(\mathcal{M}_i) p(\mathcal{D} | \mathcal{M}_i).$$

Posterior	Prior	Model evidence or marginal likelihood
-----------	-------	--

Bayes Factor: ratio of evidence for two models

$$\frac{p(\mathcal{D} | \mathcal{M}_i)}{p(\mathcal{D} | \mathcal{M}_j)}$$

Bayesian Model Comparison (2)

Having computed $p(\mathcal{M}_i|\mathcal{D})$, we can compute the predictive (mixture) distribution

$$p(t|\mathbf{x}, \mathcal{D}) = \sum_{i=1}^L p(t|\mathbf{x}, \mathcal{M}_i, \mathcal{D}) p(\mathcal{M}_i|\mathcal{D}).$$

A simpler approximation, known as *model selection*, is to use the model with the highest evidence.

Bayesian Model Comparison (3)

For a model with parameters \mathbf{w} , we get the model evidence by marginalizing over \mathbf{w}

$$p(\mathcal{D}|\mathcal{M}_i) = \int p(\mathcal{D}|\mathbf{w}, \mathcal{M}_i) p(\mathbf{w}|\mathcal{M}_i) d\mathbf{w}.$$

Note that

$$p(\mathbf{w}|\mathcal{D}, \mathcal{M}_i) = \frac{p(\mathcal{D}|\mathbf{w}, \mathcal{M}_i) p(\mathbf{w}|\mathcal{M}_i)}{p(\mathcal{D}|\mathcal{M}_i)}$$

选择最大信任的 model 来作为模型选择，而非用交叉验证，信任近似：

The Evidence Approximation (1)

The fully Bayesian predictive distribution is given by

$$p(t|\mathbf{t}) = \iiint p(t|\mathbf{w}, \beta) p(\mathbf{w}|\mathbf{t}, \alpha, \beta) p(\alpha, \beta|\mathbf{t}) d\mathbf{w} d\alpha d\beta$$

but this integral is intractable. Approximate with

$$p(t|\mathbf{t}) \simeq p(t|\mathbf{t}, \hat{\alpha}, \hat{\beta}) = \int p(t|\mathbf{w}, \hat{\beta}) p(\mathbf{w}|\mathbf{t}, \hat{\alpha}, \hat{\beta}) d\mathbf{w}$$

where $(\hat{\alpha}, \hat{\beta})$ is the mode of $p(\alpha, \beta|\mathbf{t})$, which is assumed to be sharply peaked; a.k.a. *empirical Bayes*, *type II* or *generalized maximum likelihood*, or *evidence approximation*.

The Evidence Approximation (2)

From Bayes' theorem we have

$$p(\alpha, \beta | \mathbf{t}) \propto p(\mathbf{t} | \alpha, \beta) p(\alpha, \beta)$$

and if we assume $p(\alpha, \beta)$ to be flat we see that

$$\begin{aligned} p(\alpha, \beta | \mathbf{t}) &\propto p(\mathbf{t} | \alpha, \beta) \\ &= \int p(\mathbf{t} | \mathbf{w}, \beta) p(\mathbf{w} | \alpha) d\mathbf{w}. \end{aligned}$$

General results for Gaussian integrals give

$$\ln p(\mathbf{t} | \alpha, \beta) = \frac{M}{2} \ln \alpha + \frac{N}{2} \ln \beta - E(\mathbf{m}_N) + \frac{1}{2} \ln |\mathbf{S}_N| - \frac{N}{2} \ln(2\pi).$$

Limitations of Fixed Basis Functions

- M basis function along each dimension of a D -dimensional input space requires M^D basis functions: the curse of dimensionality.
- In later chapters, we shall see how we can get away with fewer basis functions, by choosing these using the training data.

固定基存在缺陷为 NN,SVM 做铺垫, NN,SVM 都是变化基, BP 是梯度下降 error, 固定基, RBF 是聚类寻找基, SVM 是 2 次凸优化寻找基。好了,就讲到这里吧,肯定还有讲的不对,或者不足的地方,请大家一起讨论和补充,谢谢。

=====讨论=====

Wilbur_中博(1954123) 21:08:29

RBF 不是固定径向基找系数的么, SVM 也是固定基的吧, 这里寻找基是什么意思?

planktonli(1027753147) 21:09:01

SVM 是寻找那些 系数不为 0 的作为基, RBF,我说的是 RBF 神经网络, 不是 RBF 基函数, 呵呵

Wilbur_中博(1954123) 21:11:07

嗯, 但咱们现在这一章, 比如多项式基, 也可以说是寻找系数不为 0 的 x^k 吧, SVM 也仍然是固定了某一种核, 比如多项式核或者高斯核。嗯, 我知道是说 RBF 网络。

planktonli(1027753147) 21:11:40

恩,可以这么说

Wilbur_中博(1954123) 21:12:35

还有就是, 固定一组基的话, 也有很多选择, 有多项式、也有高斯、logistic 等等, 那我们应该怎么选择用什么基去做回归呢? 这一章讲得大多都是有了基以后怎么选择 w , 但怎么选择基这一点有没有什么说法。

planktonli(1027753147) 21:13:37

我说的固定指的是, SVM 不知道基是谁, 而是通过优化获取的。

Wilbur_中博(1954123) 21:13:41

或者小波傅里叶什么的。。好多基

planktonli(1027753147) 21:14:03

3.6. Limitations of Fixed Basis Functions

这里提出了固定基的问题，基的选择要看样本的几何形状，一般都是选择 gaussian，当然也可以一个个测试着弄。

Wilbur_中博(1954123) 21:15:55

SVM 里有个叫 multiple kernel learning 的，感觉像是更广泛的变化基的解决方案。嗯，就是说大多是经验性的是吧，选基这个还是蛮有趣的，我觉得。

planktonli(1027753147) 21:16:45

恩,MK 是多个 kernel 的组合，尝试用多个几何形状的 kernel 去寻找一个更 power 的。

Wilbur_中博(1954123) 21:17:05

嗯，呵呵

planktonli(1027753147) 21:17:16

恩,kernel construction 是 ML 的主要研究内容之一

Wilbur_中博(1954123) 21:18:14

好的，我没什么问题了，谢谢，以后多交流。看其他朋友还有什么问题。

planktonli(1027753147) 21:50:29

本次的讲义有些内容是群共享里的 Linear1.pdf

下次的 linear classification 主要讲的内容在群共享中为 Linear2.pdf